

NC State University
Department of Electrical and Computer Engineering
ECE 785: Spring 2019 (Dr. Dean)
Project #2:

Vectorizing the Spherical Geometry Code
by
<< SALONI SHAMBHUWANI >>

NCSU Honor Pledge: "I have neither given nor received unauthorized aid on this test or assignment."

Student's electronic signature: __Saloni Shambhuwani__
(sign by typing your name)

Course number: _____ECE-785_____

Initial code analysis

In project1 the spherical geometry code was optimized for speed with initial speed being 73-74us reducing it to 44-45us. Now to further optimize we use SIMD instructions to get parallelism on the operations that are performed in the loop. In project1 the optimized code became stagnant (on optimization) for the three loops. The loops which were major area of interest were finding closeness using fast cos approximation and removing the farthest points, second loop to find the indexes of the closest points and final loop to find the closest of all the points using cos73s and thus calculating distance and bearing. These loops still took a lot of time in finding the closeness and stalling the pipeline.

Below is the table of optimizations performed and time taken to perform them and improvement after optimization:

Step	Otimization description	Time taken for optimization	Time reduction/Speed increase compared to previous optimization code
1	Starter code		
2	Optimization1: Vectorization of cos function (fast cos (cos_mycode))	5 hrs	~50%
3	Optimization2(a): Vectorizing the array and radius comparison	0.50 hr	~20%
4	Optimization2(b): Vectorizing max value calculation and finding indexes.	7hrs	~50%
5	Optimization3: Inlining cos function	10 mins	~1%

Optimization 1: Vectorization of cos function (fast cos (cos_mycode) - initial runs to find closest waypoints)

The cos function (fast cos – $c1 + x2*(c2)$) in the project1 final code is called 164 times to find the closeness with less precision. Vectorizing this function will reduce the loop iteration by x/4. Following steps were performed to vectorize cos function:

- Inline the Calc_Closeness function into Find_Nearest_Waypoint so we can operate on four points at once.
- Changed while loop into for loop iteration of 164, incrementing it loop induction by 4
- Using vector multiplication found the values in each quadrant.
- Using vector compare and select, found the correct quadrant for each element and selected the corresponding value for that quadrant (cos values found in step iii)

- v) Unvectorized each value found from step iv and did the further operation for each lane separately.
- vi) Include `#include <arm_neon.h>` in code and Changing the Makefile to make vectorized code.
`MORE_CFLAGS = -fno-tree-loop-vectorize -fno-tree-vectorize -fno-unroll-loops`

Thus on vectorizing, I found the significant reduction in time (almost by half). The final time after optimization is given below:

Time after optimization 1: 23.42-24.17us

Below is the Perf report after optimization:

```

debian@beaglebone: ~/Project1/Project_dic
Find_Nearest_Waypoint /home/debian/Project1/Project_dic/sg
int k=0,e=0;
mov r1, s1
//for(i=0;i<163;i++)
{
    //newarr[i]=arr[closest_i]+
    if(arr[k]>=radius)
    {
        index_waypoint[e]=k;
0.00 add.w lr, sp, #96 : 0x60
ldr r5, [r4, r3]
radius= max_c * ((100-epsilon)/(100+epsilon));
vmul.f s15, s23, s15
0.03 mov r3, r5
while(strcmp(waypoints[k].Name, "END"))
b.n l07cc <Find_Nearest_Waypoint+0x1b4>
if(arr[k]>=radius)
1.06 vldmia r6!, {s14}
adds r3, #40 : 0x28
2.08 vcmp.f s15, s14
0.90 vmrs APSR_nzcv, fpscr
index_waypoint[e]=k;
29.23 itc ls
strls r1, [lr, s1, lsl #2]
//printf("Closest waypoint is %s with c as %f\n", waypoints[k].Name, arr[k]);
//memcpy(newarr[k], waypoints[i], sizeof(PT_T));
//printf("Closest waypoint is %s\n", newarr[e].Name);
e++;
1.14 adds. s1, s1, #1
}
k++;
adds r1, #1
while(strcmp(waypoints[k].Name, "END"))
0.92 ldrb r0, [r3, #14]
2.06 cmp r0, #65
bne.n l07b2 <Find_Nearest_Waypoint+0x19a>
0.23 ldrb r0, [r3, #17]
0.09 cmp r0, #78
bne.n l07b2 <Find_Nearest_Waypoint+0x19a>
0.08 ldrb r0, [r3, #18]
0.00 cmp r0, #68
bne.n l07b2 <Find_Nearest_Waypoint+0x19a>
0.01 ldrb r4, [r3, #19]
Press 'h' for help on key bindings

```

This shows that now the maximum time is spent in finding and storing the indexes in an array after fast cos. So we do the next optimization by vectorizing the array and radius comparison and finding index.

Optimization 2(a): Vectorizing the array and radius comparison

On vectorizing array and radius comparison there was few us reduction in time. Below steps were performed for vectorizing the comparison: Code is also shown below.

- i) Loading first four array elements and creating radius as the vector of 4 elements.
- ii) Using vector compare instructions to compare 4 values in one iteration
- iii) Unvectorizing compare result and finding the indexes for the one's greater than radius.

```

for(k=0;k<164;k+=4)
{
    temp_arr=vld1q_f32(&arr[k]);
    //print_float32x4(temp_arr);
    //printf("\n");
    temp_status=vcgeq_f32(temp_arr,temp_radius);
    //print_uint32x4(temp_status);
    //printf("\n");
    w1=vget_low_u32(temp_status);
    w[0]=vget_lane_u32(w1, 0);
    w[1]=vget_lane_u32(w1, 1);
    w1=vget_high_u32(temp_status);
    w[2]=vget_lane_u32(w1, 0);
    w[3]=vget_lane_u32(w1, 1);
    //printf("The values in w is %u %u %u %u",w[0],w[1],w[2],w[3]);
    for(int p=0;p<4;p++)
    {
        if(w[p]!=0)
        {
            index_waypoint[e]=k+p;
            e++;
        }
    }
}

```

This optimization/vectorization reduced the time by 6us. The time after optimization is 18.2 –18.7us.

Below perf report shows that the long time is spent in finding index and unvectorizing.

```
debian@beaglebone: ~/Project1/Project_dic
Find_Nearest_Waypoint /home/debian/Project1/Project_dic/sg
0.37      vmov.3 r3, d1[0]
          vcmpe. s13, s14
          Find_Nearest_Waypoint():
          //(r.diff_float) = c2 - max_c;
          //arr[i+2]=vget_lane_f32(v3, 0);
          arr[i+2]=c2;
          //(s.diff_float) = c3- max_c;
          //arr[i+3]=vget_lane_f32(v3, 1);
          arr[i+3]=c3;
5.88      vstr. s14, [r9, #-4]
0.34      vmrs APSR_nzcv, fpscr
7.49      it ge
          vmovge s14, s13
          vget_lane_f32():
          vmov s15, r3
          Find_Nearest_Waypoint():
          arr[i+2]=c2;
0.36      str.w r3, [r9, #-8]
          vcmpe. s14, s15
0.35      vmrs APSR_nzcv, fpscr
9.31      it ge
          vmovge s15, s14
0.34      vcmpe. s16, s15
0.29      vmrs APSR_nzcv, fpscr
7.37      it lt
          vmovlt s16, s15
          for(i=0;i<164;i+=4){
0.38      cmp r9, r4
          bne.n 1076a <Find_Nearest_Waypoint+0xc2>
          }
          // printf("max c value is %f \n",max_c);
          //radius = epsilon + acosf(arr[closest_i])*6371;
          //radius = 0.001 + arr[closest_i];
          radius= max_c * ((100-epsilon)/(100+epsilon));
0.10      vldr s15, [pc, #372] ; 10994 <Find_Nearest_Waypoint+0x2ec>
          float32_t radius_v=radius;
          int k=0,e=0;
          movs r4, #0
          {
              if(w[p]!=0)
          }
Press 'h' for help on key bindings

debian@beaglebone: ~/Project1/Project_dic
Find_Nearest_Waypoint /home/debian/Project1/Project_dic/sg
          return (float32x4_t) __builtin_neon_vmulFv4sf (__a, __b);
0.65      vld1.6 {d16-d17}, [sp, #64]
          vmul.f q6, q8, q6
          vld1q f32():
          return (float32x4_t) __builtin_neon_vld1v4sf ((const __builtin_neon_sf *) __a);
0.64      vld1.3 {d14-d15}, [r8]
          Find_Nearest_Waypoint():
          - bl cos_mycode
          add.w r8, r8, #16
          vmulq f32():
          return (float32x4_t) __builtin_neon_vmulFv4sf (__a, __b);
0.40      vldr d16, [sp, #32]
0.66      vldr d17, [sp, #40] ; 0x28
          vmul.f q9, q8, q7
          vmul.f q0, q6, q0
          vaddq f32():
          return (float32x4_t) __builtin_neon_vaddv4sf (__a, __b);
          vadd.f q0, q9, q0
          vget_lane_f32():
          return (float32_t) __builtin_neon_vget_lanev2sf (__a, __b);
0.38      vmov.3 r3, d0[0]
13.95     vmov s12, r3
0.36      vmov.3 r3, d0[1]
          Find_Nearest_Waypoint():
          //declare c1,c2,c3 and do below operations for all variables
          //v2 = vget_low_f32(temp);
          //(p.diff_float) = c - max_c;
          //arr[i]= vget_lane_f32(v2_2, 0);
          arr[i]= c;
5.91      vstr. s12, [r9, #-16]
          vget_lane_f32():
          vmov s13, r3
          vmov.3 r3, d1[1]
          vcmpe. s12, s13
          Find_Nearest_Waypoint():
          //(q.diff_float) = c1 - max_c;
          //arr[i+1]=vget_lane_f32(v2_2, 1);
          arr[i+1]=c1;
5.73      vstr. s13, [r9, #-12]
0.36      vmrs APSR_nzcv, fpscr
7.53      it ge
          Press 'h' for help on key bindings
```

Optimization 2(b): Vectorizing max value calculation and finding indexes.

In this optimization I vectorized c_max value calculation and finding the indexes of the one's which are maximum. The optimization gave significant improvement. The time reduced to almost half. Below steps were performed in this optimization:

- i) Initializing the array vector with current indexes. Initializing previous c_max vector and index vector.
- ii) Finding the c_max value using vector comparison between previous and current 4 element values of closeness.
- iii) Masking the values with one if the previous and current value is same, else updating the values with current values.
- iv) Updating the index vector for the values updated.
- v) After loop completion forwarding c_max vector and index vector for 2nd pass.

```
//vectorize c_max value calculation
//update the indexes for c_max value (selecting based on previous and new value)
//update if value has changed else take previous value
int array[]={i+0,i+1,i+2,i+3};
int32x4_t array_vector=vld1q_s32(&array[0]);
c_max_vector=vmaxq_f32(temp,c_max_vector);
uint32x4_t ene=vceqq_f32(c_max_vector_pre,c_max_vector);
//get indexes
max_value_indexes=vbslq_s32(ene,max_value_indexes,array_vector);
c_max_vector_pre=c_max_vector;
```

The code gave the reduction in time by half. Final time – 11.2-11.5us

Perf report is shown below after optimization 2(b).

```
debian@beaglebone: ~/Project1/Project_dic
Find_Nearest_Waypoint /home/debian/Project1/Project_dic/sg
0.15      cos_73(p2->Lon - p1->Lon);
        vldr s0, [r1, #12]
        pl->CosLat * p2->CosLat*
        vmul.f s18, s15, s18
        cos_73(p2->Lon - p1->Lon);
0.10      vldr s15, [sp, #72] ; 0x48
        vsub.f s0, s0, s15
        vcvt.f d0, s0
0.07      - bl cos_73
        return pl->SinLat * p2->SinLat +
0.11      vcvt.f d8, s16
        pl->CosLat * p2->CosLat*
        vcvt.f d9, s18
        return pl->SinLat * p2->SinLat +
0.14      vmla.f d8, d9, d0
0.03      vcvt.f s16, d8
        Find_Nearest_Waypoint():
        closest_i=w[t];
0.08      vcmpe. s16, s20
0.10      vmrs APSR_nzcv, fpscr
19.13      ite gt
        movgt r3, r8
0.15      movle r3, r7
        it gt
0.10      vmovgt s20, s16
        while(t<4)
        cmp.w r9, #4
        closest_i=w[t];
0.07      mov r7, r3
        while(t<4)
        beq.n 10866 <Find_Nearest_Waypoint+0x1be>
0.05      ldr.w r8, [r5, r9, lsl #2]
        b.n 10800 <Find_Nearest_Waypoint+0x158>
        // Finish calculations for The closest point

        d = acosf(max_c)*6371; // finish distance calculation
0.48      vmov.f s0, s20
        str r3, [sp, #0]
press 'h' for help on key bindings
```

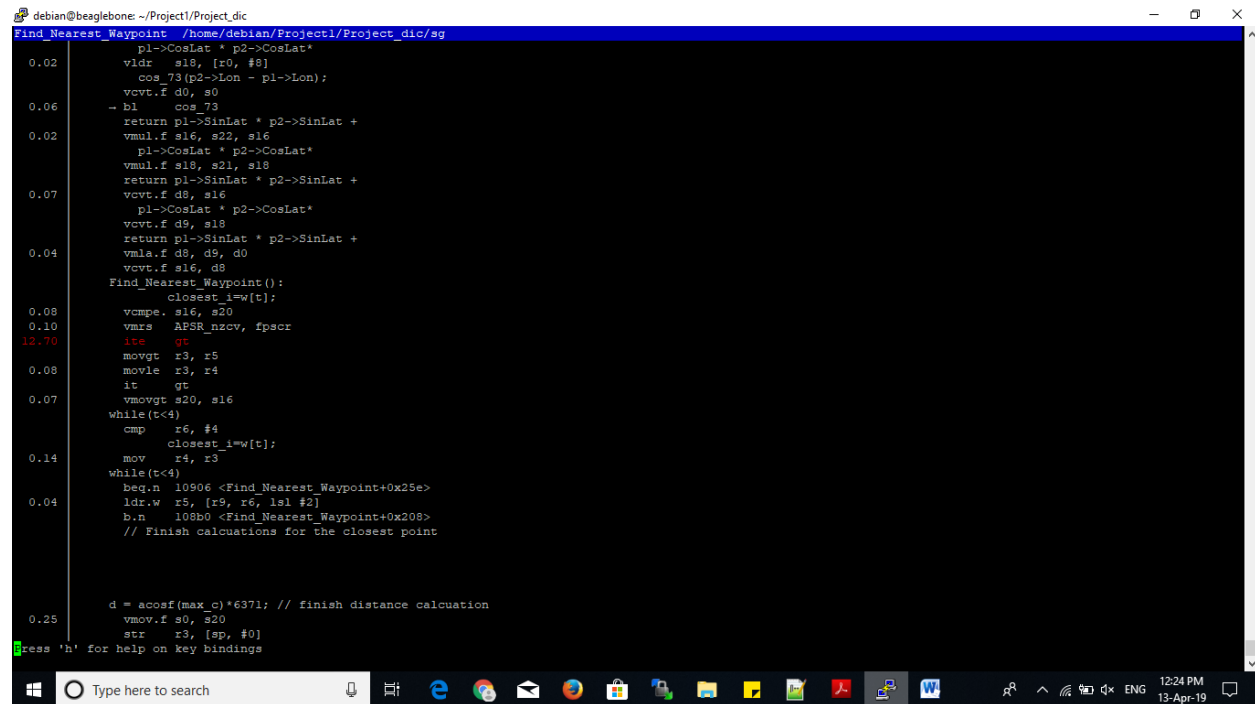
```
debian@beaglebone: ~/Project1/Project_dic
Find_Nearest_Waypoint /home/debian/Project1/Project_dic/sg
0.12      mla r4, r4, r7, r6
0.07      vldr s14, [sp, #72] ; 0x48
        mov r1, fp
        mov r0, s1
        Find_Nearest_Waypoint():
        d = acosf(max_c)*6371; // finish distance calculation
        vldr s16, [pc, #140] ; 1090c <Find_Nearest_Waypoint+0x264>
        Calc_Bearing():
        term1 = sinf(pl->Lon - p2->Lon)*p2->CosLat;
0.10      vldr s15, [r4, #12]
        Find_Nearest_Waypoint():
        d = acosf(max_c)*6371; // finish distance calculation
        vmul.f s16, s0, s16
0.08      vsub.f s0, s14, s15
        - bl sincosf
        vldr s1, [sp, #88] ; 0x58
        Calc_Bearing():
        term1 = sinf(pl->Lon - p2->Lon)*p2->CosLat;
6.10      vldr s15, [r4, #8]
        pl->SinLat*p2->CosLat*cosf(pl->Lon - p2->Lon);
0.05      vldr s13, [sp, #68] ; 0x44
        vmul.f s1, s1, s15
        vldr s0, [sp, #92] ; 0x5c
        term2 = pl->CosLat*p2->SinLat -
        vldr s14, [r4, #4]
        pl->SinLat*p2->CosLat*cosf(pl->Lon - p2->Lon);
        vmul.f s1, s1, s13
        angle = atan2f(term1, term2) * (180/PI);
0.01      vmul.f s0, s0, s15
        vldr s15, [sp, #64] ; 0x40
        vmls. s1, s15, s14
        - bl atan2f_finite
0.47      vldr s15, [pc, #80] ; 10910 <Find_Nearest_Waypoint+0x268>
        if (angle < 0.0)
0.08      ldr r3, [sp, #0]
        Find_Nearest_Waypoint():
        b = Calc_Bearing(&ref, &(waypoints[closest_i]) );
        *distance = d;
        *bearing = 360-b;
        *name = (char *) (waypoints[closest_i].Name);
        movs r2, #40 ; 0x28
press 'h' for help on key bindings
```

Optimization 3: Inlining cos function

This optimization was more of the instinct than analysis. I was calling cos calculation function (vectorized cos-fast cos) from sincos.c. On inlining the cos function with geometry.c code it gave 1us reduction in time.

Final time: 10.08-10.3us

The perf report after optimization is shown below:



```
debian@beaglebone: ~/Project1/Project_dic
Find_Nearest_Waypoint /home/debian/Project1/Project_dic/sg
0.02      p1->CosLat * p2->CosLat*
          vldr    s18, [r0, #8]
          cos_73(p2->Lon - p1->Lon);
          vcvtr.f d0, s0
0.06      - b1      cos_73
          return p1->SinLat * p2->SinLat +
0.02      vmul.f s16, s22, s16
          p1->CosLat * p2->CosLat*
          vmul.f s18, s21, s18
          return p1->SinLat * p2->SinLat +
0.07      vcvtr.f d8, s16
          p1->CosLat * p2->CosLat*
          vcvtr.f d9, s18
          return p1->SinLat * p2->SinLat +
0.04      vmul.f d8, d9, d0
          vcvtr.f s16, d8
          Find_Nearest_Waypoint():
          closest_i=w[t];
0.08      vcmpe. s16, s20
0.10      vmrs   AFSR_nzcv, fpscr
12.70     ite    qv
0.08      movgt r3, r5
          movle r3, r4
          it     gt
0.07      vmovgt s20, s16
          while(t<4)
          cmp    r6, #4
          closest_i=w[t];
0.14      mov    r4, r3
          while(t<4)
          beq.n 10906 <Find_Nearest_Waypoint+0x25e>
          ldr.w r5, [r3, r6, lsl #2]
          b.n 108b0 <Find_Nearest_Waypoint+0x208>
          // Finish calculations for the closest point

          d = acosf(max_c)*6371; // finish distance calculation
0.25      vmov.f s0, s20
          str    r3, [sp, #0]
press 'h' for help on key bindings
```

There is still time used up in the loop unvectorizing the indexes for 2nd pass. But second pass is double precision calculation. So the single precision vectorization could not be performed.

The vectorization combined with project1 optimization reduced the time from 75us to 10us which is the significant improvement in the speed of the code.

Lessons learned:

Optimisations get more and more difficult as time progresses as we need to find the inefficiencies at deeper and deeper level. Initially we can just look at the high-level code and figure out the optimizable sections, then we need to research a little bit more and get into the intricacies of the language. But, after that one needs to focus on the object code and see how the hardware/system behaves in presence of certain instructions. The programmer also needs to alter certain pre-made functions to better fit the application requirements thereby performing efficiently. In case of this code, it was already optimised to a very higher degree which made finding inefficiencies difficult.